

MPTE: Theory-Validation Simulation

Results report (internal, for discussion)

July 6, 2026

1 What this validates

The inferential theory (Theorems 1–3) is derived for the *linear* MPTE estimator: PCA on the attention-weighted panel $\tilde{Z} = BZA_z$ with fixed operators B , A_z . This report validates that estimator on a controlled union-factor linear DGP with known ground truth, and bridges it to the nonlinear model. Four experiments: E1 consistency (Theorem 1), E2 coverage (Theorems 2–3), E3 efficiency from transfer learning, E4 nonlinear bridge.

Model note: recovering the two attention operators. The theory uses two operators, a temporal B ($T \times T$) that reweights time and a cross-sectional A_z ($N \times N$) that reweights variables. A single self-attention layer produces only *one* square matrix, over whatever we treat as its tokens: attend over the T time rows and we get a $T \times T$ matrix (B); attend over the N variable columns and we get an $N \times N$ matrix (A_z). One attention therefore sees one axis only. The empirical pipeline flattens the whole panel into one long token stream, runs one large attention over it, and recovers B and A_z afterward by averaging that attention over variables and over time. At the sizes needed to trace asymptotic rates the flattening is infeasible: a flattened $T \times N$ panel has $T \cdot N$ tokens and a $(T \cdot N)^2$ attention, for example 160,000 tokens and 2.56×10^{10} entries at $N = T = 400$. We therefore use *axial* attention: one attention over the N columns yields A_z and one attention over the T rows yields B , directly and at cost $N^2 + T^2 \approx 3.2 \times 10^5$. This produces the same two operators the paper defines, obtained directly rather than by averaging a flattened map. The value map is fixed to the identity ($W^v = I$, so $V = Z$): the attention reweights the panel Z itself, which keeps $\tilde{Z} = BZA_z$ an approximate factor model in the same factors and loadings, so PCA on $\tilde{Z}$ recovers them.

Setup. The data-generating process is the union-factor linear model of the design note ($k = 4$ factors, $N_x:N_y = 2:1$, iid idiosyncratic noise). Every quantity averages over 2000 Monte Carlo replications: per (T, N) cell in E1 (grid $N = T \in \{80, 160, 320, 640, 1000\}$), in E2 and E3, and over 2000 out-of-sample panels in E4 (at $N = T = 300$).

2 E1 — Consistency (Theorem 1)

Theorem 1 gives consistency of the factor, loading, and common-component estimates. We track the common component $C = F\Lambda^\top$, the rotation-free combination of the factor and loading estimates: it needs no rotation alignment and summarizes both in one quantity. We report its relative error $\|\hat{C} - C\|_F^2 / \|C\|_F^2$ as the panel grows along $N = T$. The relative form is scale-invariant, so it is comparable across operator arms, whereas the absolute error confounds estimation quality with the operator-dependent scale of the target $C = BCA_z$. Consistency requires the relative error to decline toward zero, and a log-log slope near -1 matches the classical PCA rate $1/T + 2/N$.

We report three operator arms, each grounded in the paper. The oracle arm fixes $A_z = I$ and $B = I$, the pure-theory reference. The parameter-free arm uses the paper’s simplest attention, $Q = K = Z$, so that $A_z = \text{softmax}(Z^\top Z/\cdot)$ carries no learned projection. The learned arm reads the frozen, scaled attention off the trained linear model, estimated on a separate panel and then held fixed, the sample-splitting convention of the paper’s Remark 1. All three arms converge at a slope near one (Figure 1, Table 1), so they agree, which is the property the design asks for: the data-driven operators inherit Theorem 1 at the same rate as the oracle.

A note on Assumption A.7. For the two data-driven arms the cross-sectional operator norm $\|A_z\|_{\text{op}}$ grows with N rather than staying bounded, so A.7 as stated does not hold along this path. The convergence rate is nevertheless unaffected, as panel (d) read against panels (a)–(c) shows. We read this as evidence that A.7 is sufficient but stronger than necessary in these designs. Whether it can be relaxed to a slowly growing operator norm, with the rate re-derived, is a question for the theory that we leave open here.




Figure 1: E1: relative estimation error against $N = T$ on log-log axes for (a) the common component C , (b) the factors F , and (c) the loadings Λ , each with a slope near -1 for all three operator arms; (d) the cross-sectional operator norm $\|A_z\|_{\text{op}}$, which grows for the data-driven arms while the rates in (a)–(c) hold.

Operator arm	slope C	slope F	slope Λ	$\ A_z\ _{\text{op}}$ at N_{max}
oracle ($A_z=I$)	-1.03	-1.01	-1.10	1.0
parameter-free ($Q=K=Z$)	-0.86	-0.69	-1.53	5.6
learned, frozen	-1.08	-1.30	-0.85	13.8

Table 1: E1: fitted slope of the relative error against N (target -1) for the common component, factors, and loadings, per operator arm, with the operator norm at the largest N .

3 E2 — Coverage / asymptotic normality (Theorems 2–3)

Empirical coverage of CIs for the Y -strong common component $C_{y,i,t}$, using the analytic plug-in variance from the paper. For the baseline iid DGP the long-run covariances collapse to $\Omega = \sigma^2 \Sigma_{F,y}$, $\Xi = \sigma^2 \Sigma_{\Lambda,y_s}$, and we studentize with the general Theorem-3(iii) variance $\text{Var}(\hat{C} - C) = \sigma_F^2/N + \sigma_\Lambda^2/T$ (the F - and Λ -dominant cases are its limits). Coverage is near nominal in all three regimes (Table 2) and the studentized statistic is standard normal (Figure 2).

Regime	N	T	Cov. 90%	Cov. 95%	z sd
F-dominant	50	400	0.878	0.936	1.057
Lambda-dominant	400	50	0.886	0.941	1.021
mixed	200	200	0.894	0.947	1.006

Table 2: E2: empirical coverage of 90%/95% CIs and the studentized-statistic sd, by regime.




Figure 2: E2: QQ plot of the studentized $C_{y,i,t}$ against $\mathcal{N}(0, 1)$ (mixed regime).

4 E3 — Efficiency from transfer learning

The joint (X, Y) estimator of $C_{y,i,t}$ is more precise than the Y -only estimator; the gain grows with the auxiliary size N_x and saturates because only one factor is shared (Figure 3, Table 3).




Figure 3: E3: MSE ratio (Y-only / joint) vs. N_x .

N_x	MSE joint	MSE Y-only	ratio
25	0.1119	0.1016	0.908
50	0.1032	0.1018	0.986
100	0.0955	0.1015	1.062
200	0.0905	0.1018	1.125
400	0.0884	0.1015	1.149

Table 3: E3: joint vs. Y-only common-component MSE and their ratio.

5 E4 — Nonlinear bridge

On linear-truth data we ask whether the nonlinear model builds the factor representation the theory describes (Figure 4). We measure two subspaces against the true attended factors $F^{(B)} = BF$ by canonical correlation: the nonlinear model’s k -dim latent, and the linear PCA estimator’s factors. The linear estimator recovers the full attended factor space. The nonlinear model, trained on a single-target forecast objective with a deliberately minimal architecture, recovers the target-relevant direction strongly (leading canonical correlation ≈ 0.98) while leaving the orthogonal factors to the linear estimator, since a single-target objective only needs the direction that predicts its target. The forecast-accuracy equivalence between the nonlinear model and the linear ablation is the paper’s

Table 1 result; here we add the factor-space statement against known ground truth for the target signal.




Figure 4: E4: canonical correlations with $F^{(B)}$, nonlinear latent vs. linear PCA. The nonlinear model recovers the leading, target-relevant factor; the linear estimator recovers the full space.

6 Learned operators

Figure 5 shows the frozen cross-sectional operator A_z and temporal operator B read off the trained linear model on one panel. The attention concentrates on a few variables and time points, the bright stripes, which is the structure behind the operator norm that grows with N in Figure 1(d).




Figure 5: Learned, frozen A_z ($N \times N$) and B ($T \times T$) from the independent-sample linear model.

Addition after July 2nd

The three exhibits below were added after the first read, following the July email. Each answers one request directly, and each is reported as it comes out, including the two where the clean picture does not appear. We think the honest version is the informative one.

A1. E1: error against the theoretical rate $\bar{\alpha}$

The request was to plot the factor and loading errors against the theoretical rate $\bar{\alpha}$ rather than against N . Figure 6 shows the relative errors on log-log axes against $\bar{\alpha} = 1/T + 2/N$. For the oracle the two move together on the slope-1 line, exactly as Theorem 1 asks. For the two data-driven arms the points scatter off the line: their realized $\bar{\alpha}$ absorbs the cross-sectional operator norm $\|A_z\|_{\text{op}}$, which grows with N and is non-monotone across the grid, so $\bar{\alpha}$ is no longer the clean scalar the plot needs on its horizontal axis. This is the same fact the main-text A.7 note reports, seen from the rate side. The consistency claim itself is untouched: the error-against- N panels (Figure 1) stay on slope -1 for every arm, so the data-driven operators do converge at the full rate. We keep both views and read them together.




Figure 6: Addition: relative error against the theoretical rate $\bar{\alpha}$ on log-log axes for (a) the factors and (b) the loadings, one series per operator arm, with a slope-1 reference. The oracle follows the reference; the data-driven arms scatter because their realized $\bar{\alpha}$ carries the growing operator norm.

A2. E4: block-wise factor recovery

The request was to split the E4 canonical correlations by factor block: the Y -strong block $F_S^{(B)}$, which the target loads on, against the rest block $F_R^{(B)}$ of X -only factors. Table 4 reports both blocks for the nonlinear latent and for the linear PCA factors. The linear estimator recovers both blocks near one, as a full-panel estimator should. The nonlinear latent recovers the leading Y -strong direction strongly ($\text{cc}_1 \approx 0.98$) but not the second Y -strong direction, and it also correlates with the rest block. The reason is that the latent is a linear read-out of the whole attended panel $\tilde{Z}$, and $\tilde{Z}$ carries the X columns, so the rest factors sit inside the representation regardless of its dimension. We also tried shrinking the latent to $k = 2$; that made the Y -strong recovery worse, not sharper, and did not separate the blocks. So the defensible reading of E4 is a statement about the target-relevant *direction*

the forecast objective identifies, not about isolating a full Y -strong *block*. A clean block statement would need a different probe, the subspace the forecast head actually uses, or a design where the target loads on the whole block.

Model	Y-strong block $F_S^{(B)}$		rest block $F_R^{(B)}$	
	cc1	cc2	cc1	cc2
nonlinear latent	0.980	0.585	0.922	0.616
linear PCA	0.991	0.896	0.990	0.970

Table 4: Addition: block-wise canonical correlations with the true attended factors, the Y -strong block $F_S^{(B)}$ against the rest block $F_R^{(B)}$, for the nonlinear latent and the linear PCA factors ($k = 4$, $N = T = 300$).

A3. E2: coverage under the learned operators

The main-text E2 uses the oracle operators, where the assumptions hold exactly. The request was to add the learned, frozen operators alongside. Table 5 reports both. The oracle stays near nominal in all three regimes; the learned arm under-covers sharply, with the studentized statistic spreading well past unit variance. The cause is specific: the analytic plug-in uses the iid collapse $\Omega = \sigma^2 \Sigma_{F,y}$, $\Xi = \sigma^2 \Sigma_{\Lambda, y_s}$, but once the panel is passed through a concentrated A_z the attended noise $B e A_z$ is no longer iid, so that collapse is misspecified and the interval is the wrong width. This is not a bug to tune away: it is the finite-sample counterpart of the general Theorem-3 variance with the full Ω , Ξ , and it is the piece that pairs with the A.7 relaxation. Reporting both arms makes the boundary explicit, the theory is calibrated where its assumptions hold and needs the general variance where the operators concentrate.

Regime	N	T	oracle ($A_z=I$)			learned, frozen		
			Cov. 90	Cov. 95	z sd	Cov. 90	Cov. 95	z sd
F-dominant	50	400	0.878	0.936	1.06	0.204	0.241	6.05
Lambda-dominant	400	50	0.886	0.941	1.02	0.447	0.515	2.97
mixed	200	200	0.894	0.947	1.01	0.141	0.170	5.50

Table 5: Addition: empirical coverage of the 90%/95% CIs and the studentized-statistic sd, oracle operators against learned frozen operators, by regime. The oracle 95% cells are highlighted; the learned arm’s operator norm is $\|A_z\|_{\text{op}} \approx 4.5$ (F-dominant), 10.1 (Λ -dominant), 8.8 (mixed).

Addition after July 6th

This round follows the four items of the July email, in order. The first two are settled and we adopt the reading proposed there. The third is the coverage work, where we ran all three checks. The fourth is the projected operator arm, which we correct into a blend and, once it points the way, close the coverage story with a boundary exhibit in the paper’s own quantities.

B1. E1 axis: the realized rate $\bar{\alpha}$ per arm

The horizontal axis of Figure 6 is each arm’s own realized three-term rate $\bar{\alpha}$, computed from that arm’s operators; it equals $1/T + 2/N$ only at the oracle, where $A_z = I$ and $B = I$. Read this way, Theorem 1 is an upper bound, error = $O_P(\bar{\alpha})$, so points on or below the slope-1 line are what it permits. The bound is sharp at the oracle, where A.7 holds, and conservative under the concentrated data-driven operators, where the slack sits in the middle term of $\bar{\alpha}$.

B2. E4 blocks: a one-dimensional identified set

With a single-target forecast objective the identified set is one-dimensional: the forecast needs only the Λ_i combination of the strong factors, so a leading canonical correlation near 0.98 on one Y -strong direction, and no more, is what the theory predicts. The pickup on the rest block is something the theory neither requires nor forbids. This is the statement we attach to Table 4.

B3. Coverage under the learned operators

We ran the three checks. Table 6 collects coverage for the oracle and four learned variants (raw, block-restricted, clipped, blend) under three variance channels: the iid plug-in, the feasible general plug-in, and the Monte Carlo sd, alongside the A.7 quantities $\|A_z\|_{\text{op}}$ and PR/N (the participation ratio $\text{PR} = \text{tr}(A_z^\top A_z)^2 / \|A_z^\top A_z\|_F^2$, the effective dimension; $\text{tr}/N = 1$ for all by the scaling).

(a) *Re-studentize with the Monte Carlo variance.* Dividing the same learned-arm draws by the sd measured in simulation restores the distributional shape, the studentized statistic is close to normal with the QQ bulk on the line (Figure 7) and coverage near nominal in the F- and Λ -dominant regimes, but a location term remains: coverage still falls short, unevenly across regimes (Table 6, MC column). So the width was part of the problem; a residual bias in the common component is the rest.

(b) *Block restriction.* The E2 operators were not block-restricted, so we reran with the X-to-Y cross terms zeroed. The collapse under the iid plug-in barely moves (Table 6, learned rows), so it is not identification leakage from the X block; it is the interval itself.

(c) *The feasible general plug-in.* The attended noise $\tilde{e} = BeA_z$ has a known covariance up to one scalar, $\text{Cov}(\tilde{e}_{ti}, \tilde{e}_{sj}) = \sigma^2 (BB^\top)_{ts} (A_z^\top A_z)_{ij}$, so we estimate σ^2 from the PCA residuals and rebuild the two score covariances with $(A_z^\top A_z)$ and $(BB^\top)$ inserted where the iid collapse used identities. This is the Theorem-3 variance specialized to our frozen-operator design. It reduces to the iid plug-in at the oracle exactly, and it recovers most of the width the iid collapse discards, but on the raw and block-restricted arms it stops short of nominal.

The reason it stops short is a finite-sample bias that does not close as the panel grows, and it is a statement in the paper’s own quantities. Figure 8 sweeps $N = T$ for the learned arm and reports, averaged over operator seeds and cells, the standardized bias and the coverage under the Monte Carlo and general widths, next to $\sqrt{N} \bar{\alpha}$. The middle term of $\bar{\alpha}$ equals $1/\text{PR}$, so a bounded participation ratio makes $\bar{\alpha}$ stall and $\sqrt{N} \bar{\alpha}$ grow rather than vanish; the Theorem 2–3 growth conditions $\sqrt{T} \bar{\alpha} \rightarrow 0$ and $\sqrt{N_{\text{eff}}} \bar{\alpha} \rightarrow 0$ therefore fail, and the $o_p(1)$ remainders surface as exactly the non-vanishing bias we measure. The single-target attention concentrates on a bounded set of hub variables, so PR/N keeps shrinking and the effective dimension does not grow with N . This does not contradict E1: the

Frobenius-norm error in Figure 1 is carried by the high-signal hub cells the attention favors, while the Y -block target cells keep a small relative bias.

Regime	Operator arm	$\ A_z\ _{\text{op}}$	PR/ N	iid	general	MC
F-dominant	oracle ($A_z=I$)	1.0	1.000	0.936	0.936	–
F-dominant	learned, raw	4.5	0.104	0.241	0.879	0.949
F-dominant	learned, block-restricted	3.8	0.140	0.285	0.860	0.949
F-dominant	learned, clipped	3.1	0.193	0.369	0.878	0.939
F-dominant	learned, blend	2.7	0.334	0.542	0.932	0.954
Lambda-dominant	oracle ($A_z=I$)	1.0	1.000	0.941	0.941	–
Lambda-dominant	learned, raw	10.1	0.029	0.515	0.825	0.939
Lambda-dominant	learned, block-restricted	8.7	0.035	0.606	0.787	0.922
Lambda-dominant	learned, clipped	3.4	0.133	0.790	0.898	0.950
Lambda-dominant	learned, blend	3.1	0.363	0.834	0.930	0.948
mixed	oracle ($A_z=I$)	1.0	1.000	0.947	0.947	–
mixed	learned, raw	8.8	0.028	0.170	0.433	0.697
mixed	learned, block-restricted	8.5	0.032	0.132	0.405	0.643
mixed	learned, clipped	3.4	0.120	0.482	0.908	0.940
mixed	learned, blend	2.8	0.359	0.718	0.935	0.945

Table 6: Addition: E2 coverage of the 95% CI by regime and operator arm, with the A.7 quantities $\|A_z\|_{\text{op}}$ and PR/ N , under the iid plug-in, the feasible general plug-in, and the Monte Carlo sd. The clip and blend arms apply the same projection to the temporal operator B ; the general-plug-in cells for the arms inside the theory’s domain (oracle and blend) are highlighted.

E2: learned-arm studentized statistic (mixed regime)




Figure 7: Addition: QQ of the learned-arm studentized statistic (mixed regime) under the iid plug-in, the general plug-in, and the Monte Carlo sd. The iid tails fan out; the general and Monte Carlo channels pull the bulk onto the line but leave a residual off-center from the finite-sample bias.

B4. The projected arms: clip and blend

A.7’s two halves are jointly an anti-concentration condition. If $\|A_z\|_{\text{op}} \leq \kappa$ and $\text{tr}(A_z^\top A_z)/N = c_A$, then $\text{PR}/N \geq c_A/\kappa^2$; each half alone is satisfiable by a concentrated operator, and the two failing arms show it. The raw arm keeps the trace half by scaling but its op-norm grows, so it fails the norm half. We then try two projections that are deterministic functions of the frozen operator, so the sample-splitting convention is intact, and we apply each to *both* operators, since the temporal B is concentrated on the same footing as A_z and drives the loading-CLT term.




Figure 8: Addition: (a) coverage against $N = T$ under the Monte Carlo and general widths, both plateauing below the nominal 0.95; (b) the standardized bias and the growth product $\sqrt{N} \bar{\alpha}$, neither of which vanishes, because the learned operator stays concentrated (its participation ratio does not grow with N).

The *clip* caps the singular values at $\kappa = 3$ and re-applies the trace scaling. For an effectively low-rank spectrum this has no fixed point: clipping removes top mass but cannot create the tail mass the trace half needs, so the rescale nudges the op-norm back above κ and the arm lands on the A.7 boundary, PR/N just above c_A/κ^2 (Table 6). Coverage improves but stops short. The *blend* adds the missing tail mass: clip at $\kappa_0 = 3$, add δI with $\delta = 1$, then trace-rescale once. It is a shrinkage of the learned attention toward the identity, giving every variable a baseline self-weight, and it lands comfortably inside the domain, $PR/N \approx 0.35$ with the op-norm at or below 3. Under the general plug-in the blend arm is near nominal in all three regimes (Table 6, blend rows), and its Monte Carlo coverage is nominal, so it validates the domain where the clip only reaches its edge.

Fixing both axes is necessary, and which axis binds is regime-dependent in exactly the way the two CLT terms predict. Blending A_z alone nearly closes the F-dominant regime, where small N makes the factor-CLT term (carrying A_z) dominant, but not the Λ -dominant regime; blending B alone does the reverse, closing Λ -dominant, where small T makes the loading-CLT term (carrying B) dominant; only blending both closes the mixed regime. The binding quantity is the pair of anti-concentration conditions, one per axis.

B5. The boundary in theorem-native units

Figure 9 sweeps the blend shrinkage δ from 0, the concentrated end where the blend reduces to the clip, toward 5, where it converges to the identity, applied to both operators, and plots the general-plug-in coverage against the realized effective rank PR/N (equivalently the middle term of $\bar{\alpha}$, since it equals $1/PR$). As δ lifts PR/N through the A.7 threshold c_A/κ_0^2 , coverage climbs to nominal: inference switches on precisely as the operators re-enter the theory’s domain. The boundary is a single curve in the paper’s own quantity.




Figure 9: Addition: general-plug-in coverage of the blend arm against the effective rank PR/N as the shrinkage δ grows from 0 to 5 (both operators blended), one curve per regime, with the A.7 threshold c_A/κ_0^2 marked. Coverage reaches nominal as the effective rank is restored.